

Application of Deep Learning to Text and Image Data

Module 2, Lab 1: Processing Text

In this notebook, you will learn techniques to analyze and process text data. Text processing is known as *natural language processing (NLP)* and is an important topic because of how much information is communicated through text. Knowing how to handle text will help you build models that perform better and are more useful.

You will learn the following:

- What a word cloud is and how to create one
 - How to use stemming and lemmatization
 - What part-of-speech tagging is and how it impacts text processing
 - How to use named entity recognition to sort data
-

You will be presented with two kinds of exercises throughout the notebook: activities and challenges.

No coding is needed for an activity. You try to understand a concept, answer questions, or run a code cell.

Challenges are where you can practice your coding skills.

Index

- [Word cloud](#)
 - [Part-of-speech tagging](#)
 - [Stemming and lemmatization](#)
 - [Named entity recognition](#)
-

Initial Setup

First let's put everything in place.

```
!pip install -U -q -r requirements.txt
```

Install the [spaCy](#) library. This will be used to perform some NLP tasks in the lab.

```
!python -m spacy download en_core_web_sm
```

```
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages/torch/cuda/__init__.py:551: UserWarning: Can't initialize NVML
```

```
warnings.warn("Can't initialize NVML")
```

```
Collecting en-core-web-sm==3.7.1
```

```
  Downloading
```

```
https://github.com/explosion/spacy-models/releases/download/en_core_web_sm-3.7.1/en_core_web_sm-3.7.1-py3-none-any.whl (12.8 MB)
```

```
12.8/12.8 MB 74.5 MB/s eta
```

```
0:00:0000:0100:01
```

```
Requirement already satisfied: spacy<3.8.0,>=3.7.2 in
```

```
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from en-core-web-sm==3.7.1) (3.7.2)
```

```
Requirement already satisfied: spacy-legacy<3.1.0,>=3.0.11 in
```

```
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.0.12)
```

```
Requirement already satisfied: spacy-loggers<2.0.0,>=1.0.0 in
```

```
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.0.5)
```

```
Requirement already satisfied: murmurhash<1.1.0,>=0.28.0 in /home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.0.12)
```

```
Requirement already satisfied: cymem<2.1.0,>=2.0.2 in
```

```
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2.0.11)
```

```
Requirement already satisfied: preshed<3.1.0,>=3.0.2 in /home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.0.9)
```

```
Requirement already satisfied: thinc<8.3.0,>=8.1.8 in
```

```
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (8.2.5)
```

```
Requirement already satisfied: wasabi<1.2.0,>=0.9.1 in /home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.1.3)
```

```
Requirement already satisfied: srsly<3.0.0,>=2.4.3 in
```

```
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2.5.1)
```

```
Requirement already satisfied: catalogue<2.1.0,>=2.0.6 in /home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2.0.10)
```

```
Requirement already satisfied: weasel<0.4.0,>=0.1.0 in /home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (0.3.4)
```

```
Requirement already satisfied: typer<0.10.0,>=0.3.0 in /home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (0.9.4)
```

```
Requirement already satisfied: smart-open<7.0.0,>=5.2.1 in /home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from
```

spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (6.4.0)
Requirement already satisfied: tqdm<5.0.0,>=4.38.0 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (4.66.1)
Requirement already satisfied: requests<3.0.0,>=2.13.0 in /home/ec2-
user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from
spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2.31.0)
Requirement already satisfied: pydantic!=1.8,!1.8.1,<3.0.0,>=1.7.4 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.10.13)
Requirement already satisfied: jinja2 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.1.2)
Requirement already satisfied: setuptools in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (68.2.2)
Requirement already satisfied: packaging>=20.0 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (21.3)
Requirement already satisfied: langcodes<4.0.0,>=3.2.0 in /home/ec2-
user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from
spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.5.0)
Requirement already satisfied: numpy>=1.19.0 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.23.5)
Requirement already satisfied: language-data>=1.2 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from langcodes<4.0.0,>=3.2.0->spacy<3.8.0,>=3.7.2->en-core-
web-sm==3.7.1) (1.3.0)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /home/ec2-
user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from
packaging>=20.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.0.9)
Requirement already satisfied: typing-extensions>=4.2.0 in /home/ec2-
user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from
pydantic!=1.8,!1.8.1,<3.0.0,>=1.7.4->spacy<3.8.0,>=3.7.2->en-core-
web-sm==3.7.1) (4.8.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /home/ec2-
user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from
requests<3.0.0,>=2.13.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1)
(3.3.1)
Requirement already satisfied: idna<4,>=2.5 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from requests<3.0.0,>=2.13.0->spacy<3.8.0,>=3.7.2->en-core-
web-sm==3.7.1) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from requests<3.0.0,>=2.13.0->spacy<3.8.0,>=3.7.2->en-core-
web-sm==3.7.1) (1.26.18)
Requirement already satisfied: certifi>=2017.4.17 in

```

/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from requests<3.0.0,>=2.13.0->spacy<3.8.0,>=3.7.2->en-core-
web-sm==3.7.1) (2023.7.22)
Requirement already satisfied: blis<0.8.0,>=0.7.8 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from thinc<8.3.0,>=8.1.8->spacy<3.8.0,>=3.7.2->en-core-web-
sm==3.7.1) (0.7.11)
Requirement already satisfied: confection<1.0.0,>=0.0.1 in /home/ec2-
user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from
thinc<8.3.0,>=8.1.8->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1)
(0.1.5)
Requirement already satisfied: click<9.0.0,>=7.1.1 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from typer<0.10.0,>=0.3.0->spacy<3.8.0,>=3.7.2->en-core-web-
sm==3.7.1) (8.1.7)
Requirement already satisfied: cloudpathlib<0.17.0,>=0.7.0 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from weasel<0.4.0,>=0.1.0->spacy<3.8.0,>=3.7.2->en-core-web-
sm==3.7.1) (0.16.0)
Requirement already satisfied: MarkupSafe>=2.0 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from jinja2->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1)
(2.1.3)
Requirement already satisfied: marisa-trie>=1.1.0 in
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages (from language-data>=1.2->langcodes<4.0.0,>=3.2.0-
>spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.2.1)
✓ Download and installation successful
You can now load the package via spacy.load('en_core_web_sm')

```

Import the dependencies

```

from wordcloud import WordCloud, STOPWORDS
import matplotlib.pyplot as plt
import re, string
import Stemmer
import spacy
from spacy import displacy
import pandas as pd

```

```

Matplotlib is building the font cache; this may take a moment.
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-
packages/torch/cuda/__init__.py:551: UserWarning: Can't initialize
NVML

```

```

    warnings.warn("Can't initialize NVML")

```

Next, you need to create a function to preprocess text so that only real words, not special characters and numbers, are displayed.

```
# Preprocess text
def preProcessText(text):
    # Lowercase and strip leading and trailing white space
    text = text.lower().strip()

    # Remove HTML tags
    text = re.compile("<.*?>").sub("", text)

    # Remove punctuation
    text = re.compile("[%s]" % re.escape(string.punctuation)).sub(" ", text)

    # Remove extra white space
    text = re.sub("\s+", " ", text)

    # Remove numbers
    text = re.sub(r"[0-9]", "", text)

    return text
```

Word cloud

Word clouds, which are also known as *text clouds* or *tag clouds*, help you visualize text data by highlighting the important words or phrases. Word clouds convey crucial information at a glance by making commonly occurring words bigger and bolder. These clouds are commonly used to compare and contrast two pieces of text. Word clouds are also used to identify the topic of a document.

To create a word cloud, you will use [WordCloud for Python](#).

The following text is from the [What Is Natural Language Processing \(NLP\)?](#) page on [aws.amazon.com](#).

```
text = "Natural language processing (NLP) is a machine learning \
technology that gives computers the \
ability to interpret, manipulate, and comprehend human language. \
Organizations today have large volumes \
of voice and text data from various communication channels like \
emails, text messages, social media \
newsfeeds, video, audio, and more. They use NLP software to \
automatically process this data, analyze \
the intent or sentiment in the message, and respond in real time to \
human communication. \
Natural language processing (NLP) is critical to fully and efficiently \
analyze text and speech data. \
It can work through the differences in dialects, slang, and \
grammatical irregularities typical in \
day-to-day conversations. \
```


Note that the stop words were removed before the graphic was created. This is important so that words that don't impact the meaning of the text aren't overemphasized. Can you think of some examples of stop words?

In this example, you used the precompiled list of stop words that were curated by the WordCloud for Python project. You can print a list of the stop words to make sure that they cover the stop words that you expect.

```
# Show the list of stop words
", ".join(list(STOPWORDS))
```

```
"each, only, it's, about, hers, under, wasn't, here's, can, my, been,
the, i'd, themselves, all, that's, therefore, you'd, where's, also,
into, in, who's, yours, could, didn't, shouldn't, i, haven't, above,
both, we're, like, before, through, what's, as, if, your, shall, we'd,
k, ours, than, nor, were, ever, having, r, why, can't, how, when's,
couldn't, but, here, they, mustn't, doing, weren't, let's, such, over,
on, has, why's, hadn't, don't, they're, down, are, him, once, herself,
did, he, you, do, wouldn't, against, have, you're, won't, further,
during, ought, she'll, because, until, yourself, other, who, doesn't,
of, they'll, just, myself, be, her, that, his, i've, shan't, should,
we've, where, off, then, from, while, they've, there's, we, after,
otherwise, itself, however, not, cannot, was, aren't, more, by, since,
or, which, being, isn't, again, any, does, hasn't, when, www, their,
our, you'll, you've, am, else, she'd, it, no, up, she, me, some, too,
whom, yourselves, below, get, very, most, them, between, this, own,
with, few, these, we'll, had, he'll, they'd, an, http, out, com, he's,
i'm, and, its, so, is, for, himself, she's, how's, i'll, a, theirs,
would, he'd, same, there, to, ourselves, those, what, at, hence"
```

Part-of-speech tagging

The process of classifying words into their corresponding part of speech based on definition and context is called *part-of-speech tagging*, which is also known as *POS tagging*. A part-of-speech tagger processes a sequence of words and attaches a part-of-speech tag to each word.

For this lab, you will use the the Natural Language Toolkit [spaCy](#). The **nlp()** function creates different token attributes, among them the one representing the token tag: **token.pos**. For example, the following tagged token combines the word *fly* with a noun part of speech tag, *NN*:
`tagged_tok = ('fly', 'NOUN')`.

The following table provides the meanings for the tags from a list of Universal POS tags:

Tag	Meaning
ADJ	Adjective
ADV	Adposition
ADP	Adverb

Tag	Meaning
AUX	Auxiliary
CCONJ	Coordinating Conjunction
DET	Determiner
INTJ	Interjection
NOUN	Noun
NUM	Numeral
PART	Particle
PRON	Pronoun
PROPN	Proper Noun
PUNCT	Punctuation
SCONJ	Subordinating Conjunction
SYM	Symbol
VERB	Verb
X	Other

Now you can use the tagger to tag each token or word in the following text.

Important: Always remember to preprocess the text before tagging, as we have done before in this notebook.

```
# Text sample
text
```

```
'Natural language processing (NLP) is a machine learning technology
that gives computers the ability to interpret, manipulate, and
comprehend human language. Organizations today have large volumes of
voice and text data from various communication channels like emails,
text messages, social media newsfeeds, video, audio, and more. They
use NLP software to automatically process this data, analyze the
intent or sentiment in the message, and respond in real time to human
communication. Natural language processing (NLP) is critical to fully
and efficiently analyze text and speech data. It can work through the
differences in dialects, slang, and grammatical irregularities typical
in day-to-day conversations. Companies use it for several automated
tasks, such as to: <li>Process, analyze, and archive large
documents</li> <li>Analyze customer feedback or call center
recordings</li> <li>Run chatbots for automated customer service</li>
<li>Answer who-what-when-where questions</li> <li>Classify and extract
text</li> You can also integrate NLP in customer-facing applications
to communicate more effectively with customers. For example, a chatbot
analyzes and sorts customer queries, responding automatically to
common questions and redirecting complex queries to customer support.
This automation helps reduce costs, saves agents from spending time on
redundant queries, and improves customer satisfaction.'
```



```
# Part-of-speech tagging
nlp = spacy.load("en_core_web_sm")
doc = nlp(text)
token_text = [token.orth_ for token in doc]
token_lemma = [token.lemma_ for token in doc]
token_pos = [token.pos_ for token in doc]
pd.DataFrame(zip(token_text, token_pos),
              columns=['token_text', 'token_lemma'])
```

	token_text	token_lemma
0	Natural	ADJ
1	language	NOUN
2	processing	NOUN
3	(PUNCT
4	NLP	PROPN
...
248	and	CCONJ
249	improves	VERB
250	customer	NOUN
251	satisfaction	NOUN
252	.	PUNCT

[253 rows x 2 columns]

Refer to the table in a previous cell to identify the tags that the spaCy tagger produces.

```
# uncomment the code bellow to display the dependency parse or named
# entities and their tags
#displacy.render(doc, style="dep")
```

Stemming and lemmatization

Stemming and lemmatization are two ways to process words so that a model will be more efficient. Both methods remove parts of words so that they can be grouped together.

For example, in the following sentence, "ning" would be removed from "running" so that "running" and "run" would be categorized the same.

What could make stemming and lemmatization difficult to do properly?

Stemming

Stemming is a rule-based system to convert words into their root forms by removing suffixes. This method helps to enhance similarities (if any) between sentences.

Examples:

"jumping", "jumped" -> "jump"

"cars" -> "car"

```
# let's list the language available
print(Stemmer.algorithms())

['arabic', 'armenian', 'basque', 'catalan', 'danish', 'dutch',
 'english', 'finnish', 'french', 'german', 'greek', 'hindi',
 'hungarian', 'indonesian', 'irish', 'italian', 'lithuanian', 'nepali',
 'norwegian', 'porter', 'portuguese', 'romanian', 'russian', 'serbian',
 'spanish', 'swedish', 'tamil', 'turkish', 'yiddish']
```

```
# we will use english for this example
stemmer = Stemmer.Stemmer('english')
```

```
original_text = "    This is a message to be cleaned. It may involve
some things like: <br>, ?, :, '    adjacent spaces and tabs    .    "
print(original_text)
```

```
    This is a message to be cleaned. It may involve some things like:
<br>, ?, :, '    adjacent spaces and tabs    .
```

```
# Cleaned text
cleaned_text = preProcessText(original_text)
print(cleaned_text)
```

```
this is a message to be cleaned it may involve some things like
adjacent spaces and tabs
```

```
# Use a tokenizer (nlp) and stemmer from the PyStemmer library
```

```
nlp = spacy.load("en_core_web_sm")
doc = nlp(original_text)
stemmed_sentence = []
original_sentence = []
# Tokenize the sentence
for token in doc:
    original_sentence.append(token.text)
    stemmed_sentence.append(stemmer.stemWord(token.text))
```

```
pd.DataFrame(zip(original_sentence, stemmed_sentence),
             columns=['token_text', 'token_stemmer'])
```

	token_text	token_stemmer
0		
1	This	This
2	is	is
3	a	a
4	message	messag
5	to	to
6	be	be
7	cleaned	clean
8	.	.
9	It	It

```

10      may      may
11  involve  involv
12    some    some
13  things   thing
14    like    like
15      :      :
16      <      <
17     br      br
18      >      >
19      ,      ,
20      ?      ?
21      ,      ,
22      :      :
23      '      '
24     '      '
25
26  adjacent  adjac
27    spaces  space
28     and    and
29    tabs    tab
30
31      .      .
32

```

```

stemmed_text = " ".join(stemmed_sentence)
print(stemmed_text)

```

```

This is a messag to be clean . It may involv some thing like : <
br > , ? , : , ' ' adjac space and tab .

```

From the output of the previous code cell, you can see that stemming isn't perfect. It makes mistakes, such as "messag", "involv", and "adjac". Stemming is a rule-based method that sometimes mistakenly removes suffixes from words. It does run quickly, which makes it appealing to use for massive datasets.

Lemmatization

If you aren't satisfied with the result of stemming, you can use the lemmatization instead. This method usually requires more work but gives better results.

Lemmatization needs to know the correct word position tags, such as "noun", "verb", or "adjective". You need to use another spaCy function to feed this information to the lemmatizer.

The cell below uses part of the full list of position tags listed in the previous session **Part-of-speech tagging**.

```

# get the original and the lemmatized the word/token
token_lemma = [token.lemma_ for token in doc]
token_text = [token.orth_ for token in doc]

```

```
pd.DataFrame(zip(token_text, token_lemma),
              columns=['token_text', 'token_lemma'])
```

	token_text	token_lemma
0		
1	This	this
2	is	be
3	a	a
4	message	message
5	to	to
6	be	be
7	cleaned	clean
8	.	.
9	It	it
10	may	may
11	involve	involve
12	some	some
13	things	thing
14	like	like
15	:	:
16	<	<
17	br	br
18	>	>
19	,	,
20	?	?
21	,	,
22	:	:
23	,	,
24	,	,
25		
26	adjacent	adjacent
27	spaces	space
28	and	and
29	tabs	tab
30		
31	.	.
32		

```
lemmatized_text = " ".join(token_lemma)
print(lemmatized_text)
```

```
this be a message to be clean . it may involve some thing like : <
br > , ? , : , ' ' adjacent space and tab .
```

How do the results compare? Is the lemmatized text better than the stemmed text?

Named entity recognition

Named entity recognition involves identification of key information in text and then classifying that information into predefined categories, such as person, organization, place, or date. This is one of the most popular NLP tasks.

For this section, you will use [spaCy](#). The following table lists the categories and meanings of the category labels that the spaCy module uses.

Category	Meaning
CARDINAL	Numerals that don't fall under another type
DATE	Absolute or relative dates or periods
EVENT	Named hurricanes, battles, wars, sports events, and so on
FAC	Buildings, airports, highways, bridges, and so on
GPE	Countries, cities, states
LANGUAGE	Any named language
LAW	Named documents made into laws
LOC	Non-GPE locations, mountain ranges, bodies of water
MONEY	Monetary values, including unit
NORP	Nationalities, or religious or political groups
ORDINAL	"first", "second", and so on
ORG	Companies, agencies, institutions, and so on
PERCENT	Percentage, including "%"
PERSON	People, including fictional
PRODUCT	Objects, vehicles, foods, and so on (not services)
QUANTITY	Measurements, as of weight or distance
WORK_OF_ART	Titles of books, songs, and so on

The following text was retrieved from the [Amazon - The Climate Pledge](#) page of the About Amazon website.

```
# Sample text for named entity recognition
ner_text = "Amazon and Global Optimism co-founded The Climate Pledge,
\
a commitment to net-zero \
carbon by 2040."
```

```
# Load the spaCy English pipeline for named entity recognition
NER = spacy.load("en_core_web_sm")

# Tag entities in the text
for word in NER(ner_text).ents:
    print(word.text, word.label_)

Amazon ORG
Global Optimism ORG
The Climate Pledge WORK_OF_ART
2040 DATE
```

Visualizing the Tags

spaCy has a visualizer called displaCy that you can use to visualize tags. Run the following cell to see it working.

```
# Visual tag with text
displacy.render(NER(ner_text), style="ent", jupyter=True)

<IPython.core.display.HTML object>
```

As you can see, named entity recognition can help you identify different entities in text. The process isn't always correct, but it can process large sections of text faster than a human can.

Conclusion

In this lab, you practiced using text processing techniques.

Next lab

In the next lab, you will learn about the bag-of-words (BoW) method to convert text data into numerical values.